

- [1.1] Plain-text formats are decoded directly from their bytes, while PDF documents are parsed with Apache Tika.
- [1.2] Virtual folders are expressed as slashes in the blob name, and recursion can be enabled per configured source.
- [1.3] An optional extension allowlist restricts which blobs in a folder listing are turned into documents.
- [1.4] Unsupported binaries such as images and Microsoft Office documents are skipped during folder loads.
- [1.5] When a blob is named explicitly, a deliberately unsupported type is reported as an error instead of skipped.
- [1.6] Listing follows the NextMarker pagination cursor so that large containers are fully enumerated.
- [1.7] Each returned document carries metadata: the file name, MIME type, size, and any available timestamps.
- [1.8] This fixture exists to exercise the loader end to end against a real storage account.
- [1.9] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [1.10] Text extraction should preserve the readable words on every page of a multi-page document.
- [1.11] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.

- [2.1] Virtual folders are expressed as slashes in the blob name, and recursion can be enabled per configured source.
- [2.2] An optional extension allowlist restricts which blobs in a folder listing are turned into documents.
- [2.3] Unsupported binaries such as images and Microsoft Office documents are skipped during folder loads.
- [2.4] When a blob is named explicitly, a deliberately unsupported type is reported as an error instead of skipped.
- [2.5] Listing follows the NextMarker pagination cursor so that large containers are fully enumerated.
- [2.6] Each returned document carries metadata: the file name, MIME type, size, and any available timestamps.
- [2.7] This fixture exists to exercise the loader end to end against a real storage account.
- [2.8] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [2.9] Text extraction should preserve the readable words on every page of a multi-page document.
- [2.10] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [2.11] A path without an extension that matches no blob is resolved as a folder prefix automatically.

- [3.1] An optional extension allowlist restricts which blobs in a folder listing are turned into documents.
- [3.2] Unsupported binaries such as images and Microsoft Office documents are skipped during folder loads.
- [3.3] When a blob is named explicitly, a deliberately unsupported type is reported as an error instead of skipped.
- [3.4] Listing follows the NextMarker pagination cursor so that large containers are fully enumerated.
- [3.5] Each returned document carries metadata: the file name, MIME type, size, and any available timestamps.
- [3.6] This fixture exists to exercise the loader end to end against a real storage account.
- [3.7] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [3.8] Text extraction should preserve the readable words on every page of a multi-page document.
- [3.9] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [3.10] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [3.11] Content is intentionally verbose to make the extracted output realistic rather than trivial.

- [4.1] Unsupported binaries such as images and Microsoft Office documents are skipped during folder loads.
- [4.2] When a blob is named explicitly, a deliberately unsupported type is reported as an error instead of skipped.
- [4.3] Listing follows the NextMarker pagination cursor so that large containers are fully enumerated.
- [4.4] Each returned document carries metadata: the file name, MIME type, size, and any available timestamps.
- [4.5] This fixture exists to exercise the loader end to end against a real storage account.
- [4.6] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [4.7] Text extraction should preserve the readable words on every page of a multi-page document.
- [4.8] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [4.9] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [4.10] Content is intentionally verbose to make the extracted output realistic rather than trivial.
- [4.11] Downstream retrieval-augmented generation pipelines consume these text documents as context.

- [5.1] When a blob is named explicitly, a deliberately unsupported type is reported as an error instead of skipped.
- [5.2] Listing follows the NextMarker pagination cursor so that large containers are fully enumerated.
- [5.3] Each returned document carries metadata: the file name, MIME type, size, and any available timestamps.
- [5.4] This fixture exists to exercise the loader end to end against a real storage account.
- [5.5] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [5.6] Text extraction should preserve the readable words on every page of a multi-page document.
- [5.7] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [5.8] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [5.9] Content is intentionally verbose to make the extracted output realistic rather than trivial.
- [5.10] Downstream retrieval-augmented generation pipelines consume these text documents as context.
- [5.11] The Azure Blob Storage data loader retrieves documents from one or more containers and returns them as text documents.

- [6.1] Listing follows the NextMarker pagination cursor so that large containers are fully enumerated.
- [6.2] Each returned document carries metadata: the file name, MIME type, size, and any available timestamps.
- [6.3] This fixture exists to exercise the loader end to end against a real storage account.
- [6.4] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [6.5] Text extraction should preserve the readable words on every page of a multi-page document.
- [6.6] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [6.7] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [6.8] Content is intentionally verbose to make the extracted output realistic rather than trivial.
- [6.9] Downstream retrieval-augmented generation pipelines consume these text documents as context.
- [6.10] The Azure Blob Storage data loader retrieves documents from one or more containers and returns them as text documents.
- [6.11] Plain-text formats are decoded directly from their bytes, while PDF documents are parsed with Apache Tika.

- [7.1] Each returned document carries metadata: the file name, MIME type, size, and any available timestamps.
- [7.2] This fixture exists to exercise the loader end to end against a real storage account.
- [7.3] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [7.4] Text extraction should preserve the readable words on every page of a multi-page document.
- [7.5] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [7.6] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [7.7] Content is intentionally verbose to make the extracted output realistic rather than trivial.
- [7.8] Downstream retrieval-augmented generation pipelines consume these text documents as context.
- [7.9] The Azure Blob Storage data loader retrieves documents from one or more containers and returns them as text documents.
- [7.10] Plain-text formats are decoded directly from their bytes, while PDF documents are parsed with Apache Tika.
- [7.11] Virtual folders are expressed as slashes in the blob name, and recursion can be enabled per configured source.

- [8.1] This fixture exists to exercise the loader end to end against a real storage account.
- [8.2] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [8.3] Text extraction should preserve the readable words on every page of a multi-page document.
- [8.4] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [8.5] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [8.6] Content is intentionally verbose to make the extracted output realistic rather than trivial.
- [8.7] Downstream retrieval-augmented generation pipelines consume these text documents as context.
- [8.8] The Azure Blob Storage data loader retrieves documents from one or more containers and returns them as text.
- [8.9] Plain-text formats are decoded directly from their bytes, while PDF documents are parsed with Apache Tika.
- [8.10] Virtual folders are expressed as slashes in the blob name, and recursion can be enabled per configured source.
- [8.11] An optional extension allowlist restricts which blobs in a folder listing are turned into documents.

- [9.1] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [9.2] Text extraction should preserve the readable words on every page of a multi-page document.
- [9.3] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [9.4] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [9.5] Content is intentionally verbose to make the extracted output realistic rather than trivial.
- [9.6] Downstream retrieval-augmented generation pipelines consume these text documents as context.
- [9.7] The Azure Blob Storage data loader retrieves documents from one or more containers and returns them as text documents.
- [9.8] Plain-text formats are decoded directly from their bytes, while PDF documents are parsed with Apache Tika.
- [9.9] Virtual folders are expressed as slashes in the blob name, and recursion can be enabled per configured source.
- [9.10] An optional extension allowlist restricts which blobs in a folder listing are turned into documents.
- [9.11] Unsupported binaries such as images and Microsoft Office documents are skipped during folder loads.

- [10.1] Text extraction should preserve the readable words on every page of a multi-page document.
- [10.2] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [10.3] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [10.4] Content is intentionally verbose to make the extracted output realistic rather than trivial.
- [10.5] Downstream retrieval-augmented generation pipelines consume these text documents as context.
- [10.6] The Azure Blob Storage data loader retrieves documents from one or more containers and returns them as text documents.
- [10.7] Plain-text formats are decoded directly from their bytes, while PDF documents are parsed with Apache Tika.
- [10.8] Virtual folders are expressed as slashes in the blob name, and recursion can be enabled per configured source.
- [10.9] An optional extension allowlist restricts which blobs in a folder listing are turned into documents.
- [10.10] Unsupported binaries such as images and Microsoft Office documents are skipped during folder loads.
- [10.11] When a blob is named explicitly, a deliberately unsupported type is reported as an error instead of skipped.

- [11.1] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [11.2] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [11.3] Content is intentionally verbose to make the extracted output realistic rather than trivial.
- [11.4] Downstream retrieval-augmented generation pipelines consume these text documents as context.
- [11.5] The Azure Blob Storage data loader retrieves documents from one or more containers and returns them as text documents.
- [11.6] Plain-text formats are decoded directly from their bytes, while PDF documents are parsed with Apache Tika.
- [11.7] Virtual folders are expressed as slashes in the blob name, and recursion can be enabled per configured source.
- [11.8] An optional extension allowlist restricts which blobs in a folder listing are turned into documents.
- [11.9] Unsupported binaries such as images and Microsoft Office documents are skipped during folder loads.
- [11.10] When a blob is named explicitly, a deliberately unsupported type is reported as an error instead of skipped.
- [11.11] Listing follows the NextMarker pagination cursor so that large containers are fully enumerated.

- [12.1] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [12.2] Content is intentionally verbose to make the extracted output realistic rather than trivial.
- [12.3] Downstream retrieval-augmented generation pipelines consume these text documents as context.
- [12.4] The Azure Blob Storage data loader retrieves documents from one or more containers and returns them as text documents.
- [12.5] Plain-text formats are decoded directly from their bytes, while PDF documents are parsed with Apache Tika.
- [12.6] Virtual folders are expressed as slashes in the blob name, and recursion can be enabled per configured source.
- [12.7] An optional extension allowlist restricts which blobs in a folder listing are turned into documents.
- [12.8] Unsupported binaries such as images and Microsoft Office documents are skipped during folder loads.
- [12.9] When a blob is named explicitly, a deliberately unsupported type is reported as an error instead of skipped.
- [12.10] Listing follows the NextMarker pagination cursor so that large containers are fully enumerated.
- [12.11] Each returned document carries metadata: the file name, MIME type, size, and any available timestamps.

Final page sentinel: PDF_LAST_PAGE.