

Single-Document PDF Fixture

Mock PDF document text.

- [1.1] The Azure Blob Storage data loader retrieves documents from one or more containers and returns them as text.
- [1.2] Plain-text formats are decoded directly from their bytes, while PDF documents are parsed with Apache Tika.
- [1.3] Virtual folders are expressed as slashes in the blob name, and recursion can be enabled per configured source.
- [1.4] An optional extension allowlist restricts which blobs in a folder listing are turned into documents.
- [1.5] Unsupported binaries such as images and Microsoft Office documents are skipped during folder loads.
- [1.6] When a blob is named explicitly, a deliberately unsupported type is reported as an error instead of skipped.
- [1.7] Listing follows the NextMarker pagination cursor so that large containers are fully enumerated.
- [1.8] Each returned document carries metadata: the file name, MIME type, size, and any available timestamps.
- [1.9] This fixture exists to exercise the loader end to end against a real storage account.
- [1.10] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [1.11] Text extraction should preserve the readable words on every page of a multi-page document.
- [1.12] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.

Single-Document PDF Fixture (continued)

- [2.1] Unsupported binaries such as images and Microsoft Office documents are skipped during folder loads.
- [2.2] When a blob is named explicitly, a deliberately unsupported type is reported as an error instead of skipped.
- [2.3] Listing follows the NextMarker pagination cursor so that large containers are fully enumerated.
- [2.4] Each returned document carries metadata: the file name, MIME type, size, and any available timestamps.
- [2.5] This fixture exists to exercise the loader end to end against a real storage account.
- [2.6] It contains several paragraphs so that extraction has meaningful content to return and assert on.
- [2.7] Text extraction should preserve the readable words on every page of a multi-page document.
- [2.8] The loader treats a path with a trailing slash as a folder prefix rather than an exact blob name.
- [2.9] A path without an extension that matches no blob is resolved as a folder prefix automatically.
- [2.10] Content is intentionally verbose to make the extracted output realistic rather than trivial.
- [2.11] Downstream retrieval-augmented generation pipelines consume these text documents as context.
- [2.12] The Azure Blob Storage data loader retrieves documents from one or more containers and returns them as text documents.